

Programming Exercise 1 - Linear Regression

Xinyu_Chen

February 20, 2017

0.1 Programming Exercise 1 - Linear Regression

- warmUpExercise
- Linear regression with one variable
- Gradient Descent

```
In [1]: # %load ../../standard_import.txt
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

from sklearn.linear_model import LinearRegression
from mpl_toolkits.mplot3d import axes3d

pd.set_option('display.notebook_repr_html', False)
pd.set_option('display.max_columns', None)
pd.set_option('display.max_rows', 150)
pd.set_option('display.max_seq_items', None)

#%%config InlineBackend.figure_formats = {'pdf',}
%matplotlib inline

import seaborn as sns
sns.set_context('notebook')
sns.set_style('white')
```

warmUpExercise

```
In [2]: def warmUpExercise():
        return(np.identity(5))
```

```
In [3]: warmUpExercise()
```


```
Out[3]: array([[ 1.,  0.,  0.,  0.,  0.],
               [ 0.,  1.,  0.,  0.,  0.],
               [ 0.,  0.,  1.,  0.,  0.],
               [ 0.,  0.,  0.,  1.,  0.],
               [ 0.,  0.,  0.,  0.,  1.]])
```

0.1.1 Linear regression with one variable

```
In [4]: data = np.loadtxt('data/ex1data1.txt', delimiter=',')

X = np.c_[np.ones(data.shape[0]), data[:, 0]]
y = np.c_[data[:, 1]]

In [5]: plt.scatter(X[:, 1], y, s=30, c='r', marker='x', linewidths=1)
plt.xlim(4, 24)
plt.xlabel('Population of City in 10,000s')
plt.ylabel('Profit in $10,000s');
```



Gradient Descent

```
In [6]: def computeCost(X, y, theta=[[0],[0]]):
    m = y.size
    J = 0

    h = X.dot(theta)

    J = 1 / (2 * m) * np.sum(np.square(h - y))

    return(J)
```


```
In [7]: computeCost(X,y)
```

```
Out[7]: 32.072733877455676
```

```
In [8]: def gradientDescent(X, y, theta=[[0],[0]], alpha=0.01, num_iters=1500):  
        m = y.size  
        J_history = np.zeros(num_iters)  
  
        for iter in np.arange(num_iters):  
            h = X.dot(theta)  
            theta = theta - alpha*(1/m)*(X.T.dot(h-y))  
            J_history[iter] = computeCost(X, y, theta)  
        return(theta, J_history)
```

```
In [9]: # theta for minimized cost J  
        theta , Cost_J = gradientDescent(X, y)  
        print('theta: ',theta.ravel())  
  
        plt.plot(Cost_J)  
        plt.ylabel('Cost J')  
        plt.xlabel('Iterations');
```

```
theta: [-3.63029144  1.16636235]
```



```


In [10]: xx = np.arange(5,23)
        yy = theta[0]+theta[1]*xx

        # Plot gradient descent
        plt.scatter(X[:,1], y, s=30, c='r', marker='x', linewidths=1)
        plt.plot(xx,yy, label='Linear regression (Gradient descent)')

        # Compare with Scikit-learn Linear regression
        regr = LinearRegression()
        regr.fit(X[:,1].reshape(-1,1), y.ravel())
        plt.plot(xx, regr.intercept_+regr.coef_*xx, label='Linear regression (Scikit-learn GLM)')

        plt.xlim(4,24)
        plt.xlabel('Population of City in 10,000s')
        plt.ylabel('Profit in $10,000s')
        plt.legend(loc=4);

```



```

In [11]: # Predict profit for a city with population of 35000 and 70000
        print(theta.T.dot([1, 3.5])*10000)
        print(theta.T.dot([1, 7])*10000)

[ 4519.7678677]
[ 45342.45012945]

```

```

In [12]: # Create grid coordinates for plotting
B0 = np.linspace(-10, 10, 50)
B1 = np.linspace(-1, 4, 50)
xx, yy = np.meshgrid(B0, B1, indexing='xy')
Z = np.zeros((B0.size,B1.size))

# Calculate Z-values (Cost) based on grid of coefficients
for (i,j),v in np.ndenumerate(Z):
    Z[i,j] = computeCost(X,y, theta=[[xx[i,j]], [yy[i,j]]])

fig = plt.figure(figsize=(15,6))
ax1 = fig.add_subplot(121)
ax2 = fig.add_subplot(122, projection='3d')

# Left plot
CS = ax1.contour(xx, yy, Z, np.logspace(-2, 3, 20), cmap=plt.cm.jet)
ax1.scatter(theta[0],theta[1], c='r')

# Right plot
ax2.plot_surface(xx, yy, Z, rstride=1, cstride=1, alpha=0.6, cmap=plt.cm.)
ax2.set_zlabel('Cost')
ax2.set_zlim(Z.min(),Z.max())
ax2.view_init(elev=15, azim=230)

# settings common to both plots
for ax in fig.axes:
    ax.set_xlabel(r'$\theta_0$', fontsize=17)
    ax.set_ylabel(r'$\theta_1$', fontsize=17)

```

