

Programming Exercise 3 - Multi-class Classification and Neural Networks

March 5, 2017

0.1 Programming Exercise 3 - Multi-class Classification and Neural Networks

```
In [1]: # %load ../../standard_import.txt
import pandas as pd
import numpy as np
import matplotlib as mpl
import matplotlib.pyplot as plt

# load MATLAB files
from scipy.io import loadmat
from scipy.optimize import minimize

from sklearn.linear_model import LogisticRegression

pd.set_option('display.notebook_repr_html', False)
pd.set_option('display.max_columns', None)
pd.set_option('display.max_rows', 150)
pd.set_option('display.max_seq_items', None)

%%config InlineBackend.figure_formats = {'pdf',}
%matplotlib inline

import seaborn as sns
sns.set_context('notebook')
sns.set_style('white')
```

Load MATLAB datafiles

```
In [2]: data = loadmat('data/ex3data1.mat')
data.keys()
```

```
Out[2]: dict_keys(['__header__', '__version__', '__globals__', 'X', 'y'])
```

```
In [3]: weights = loadmat('data/ex3weights.mat')
weights.keys()
```

```
Out[3]: dict_keys(['__header__', '__version__', '__globals__', 'Theta1', 'Theta2'])
```

```
In [4]: y = data['y']
        # Add constant for intercept
        X = np.c_[np.ones((data['X'].shape[0],1)), data['X']]

        print('X: {} (with intercept)'.format(X.shape))
        print('y: {}'.format(y.shape))
```

```
X: (5000, 401) (with intercept)
y: (5000, 1)
```

```
In [5]: theta1, theta2 = weights['Theta1'], weights['Theta2']

        print('theta1: {}'.format(theta1.shape))
        print('theta2: {}'.format(theta2.shape))
```

```
theta1: (25, 401)
theta2: (10, 26)
```

```
In [6]: sample = np.random.choice(X.shape[0], 20)
        plt.imshow(X[sample,1:].reshape(-1,20).T)
        plt.axis('off');
```

```
/home/ubuntu/anaconda3/lib/python3.6/site-packages/matplotlib/font_manager.py:1297:
(prop.get_family(), self.defaultFamily[fonttext]))
```



0.1.1 Multiclass Classification

Logistic regression hypothesis

$$h_{\theta}(x) = g(\theta^T x)$$

$$g(z) = \frac{1}{1 + e^{-z}}$$

```
In [7]: def sigmoid(z):
        return(1 / (1 + np.exp(-z)))
```

Regularized Cost Function

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m [-y^{(i)} \log(h_{\theta}(x^{(i)})) - (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)}))] + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$$

Vectorized Cost Function

$$J(\theta) = \frac{1}{m} ((\log(g(X\theta))^T y + (\log(1 - g(X\theta))^T (1 - y)) + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$$

```
In [8]: def lrcostFunctionReg(theta, reg, X, y):
        m = y.size
        h = sigmoid(X.dot(theta))

        J = -1*(1/m)*(np.log(h).T.dot(y)+np.log(1-h).T.dot(1-y)) + (reg/(2*m))

        if np.isnan(J[0]):
            return(np.inf)
        return(J[0])

In [9]: def lrgradientReg(theta, reg, X,y):
        m = y.size
        h = sigmoid(X.dot(theta.reshape(-1,1)))

        grad = (1/m)*X.T.dot(h-y) + (reg/m)*np.r_[[0],theta[1:].reshape(-1,1)]

        return(grad.flatten())
```

One-vs-all Classification

```
In [10]: def oneVsAll(features, classes, n_labels, reg):
        initial_theta = np.zeros((X.shape[1],1)) # 401x1
        all_theta = np.zeros((n_labels, X.shape[1])) #10x401

        for c in np.arange(1, n_labels+1):
            res = minimize(lrcostFunctionReg, initial_theta, args=(reg, features, c),
                           jac=lrgradientReg, options={'maxiter':50})
            all_theta[c-1] = res.x
        return(all_theta)

In [11]: theta = oneVsAll(X, y, 10, 0.1)
```

One-vs-all Prediction

```
In [12]: def predictOneVsAll(all_theta, features):
         probs = sigmoid(X.dot(all_theta.T))

         # Adding one because Python uses zero based indexing for the 10 columns
         # while the 10 classes are numbered from 1 to 10.
         return(np.argmax(probs, axis=1)+1)
```

```
In [13]: pred = predictOneVsAll(theta, X)
         print('Training set accuracy: {} %'.format(np.mean(pred == y.ravel())*100))
```

Training set accuracy: 93.24 %

Multiclass Logistic Regression with scikit-learn

```
In [14]: clf = LogisticRegression(C=10, penalty='l2', solver='liblinear')
         # Scikit-learn fits intercept automatically, so we exclude first column w
         clf.fit(X[:,1:],y.ravel())
```

```
Out[14]: LogisticRegression(C=10, class_weight=None, dual=False, fit_intercept=True,
                             intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
                             penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
                             verbose=0, warm_start=False)
```

```
In [15]: pred2 = clf.predict(X[:,1:])
         print('Training set accuracy: {} %'.format(np.mean(pred2 == y.ravel())*100))
```

Training set accuracy: 96.5 %

0.1.2 Neural Networks

```
In [16]: def predict(theta_1, theta_2, features):
         z2 = theta_1.dot(features.T)
         a2 = np.c_[np.ones((data['X'].shape[0],1)), sigmoid(z2).T]

         z3 = a2.dot(theta_2.T)
         a3 = sigmoid(z3)

         return(np.argmax(a3, axis=1)+1)
```

```
In [17]: pred = predict(theta1, theta2, X)
         print('Training set accuracy: {} %'.format(np.mean(pred == y.ravel())*100))
```

Training set accuracy: 97.52 %